



E V9 | Mask fusion, efficient decoding and targeted supervision

a Mask-only context / detail reconciliation (16 channels)



Correction feeds masks only; no C4/C5-width computation at P2; no CPU edge operator.

b Replace the wide prototype stack; retain fine phase decoding



c Training-only additions, tested separately

Tiny Dice: visible area <256 px²
mean mask coefficients per GT
base BCE kept for every positive

Quality negatives: ≤32 per image
exclude all GT boxes +8 px margin
no new inference branch

Keep visible-boundary / neighbor loss
no convex fill or circle target
Copy-Paste recipe inherited